#### **Data Mining - Unsupervised Machine Learning**



## **Data Clustering**

#### **Prof. Yannis Velegrakis**

University of Trento & Utrecht University

https://velgias.github.io

## **Use Cases**

- Tourism: Understanding Points of Interest
  - Transportation
  - Shops / Services
  - Security / Safety
- Security: Understanding High Risk Areas
- Population Understanding: Types of citizens and Needs
- Health: High Risk Areas (The old story of London City)
- Software Usage in the Offices:
  - Identify best needed software packages
- Home Services (Gas, Water, Electricity) Consumption
- Social Media: What the citizens want/think
- Telecommunication Data: Identify Families

- Other types of Data:
  - Segments: Transportation Routes, Tourism Paths, Super Market Shopping Plans
  - Graphs: Communities
  - Streams: City Sensors

## **High Dimensional Data**

 Given a cloud of data points we want to understand its structure



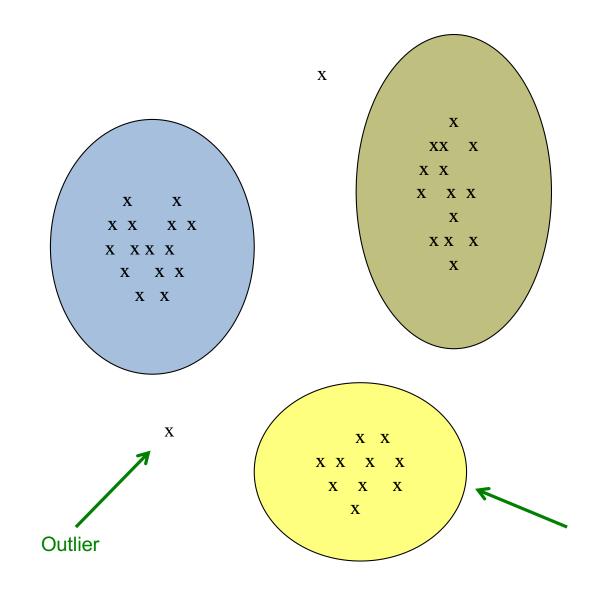
## The Problem of Clustering

- Given a set of points, with a notion of distance between points, group the points into some number of clusters, so that
  - Members of a cluster are close/similar to each other
  - Members of different clusters are dissimilar

## • Usually:

- Points are in a high-dimensional space
- Similarity is defined using a distance measure
  - ◆Euclidean, Cosine, Jaccard, edit distance, ...

# **Example: Clusters & Outliers**



# Clustering is a hard problem!

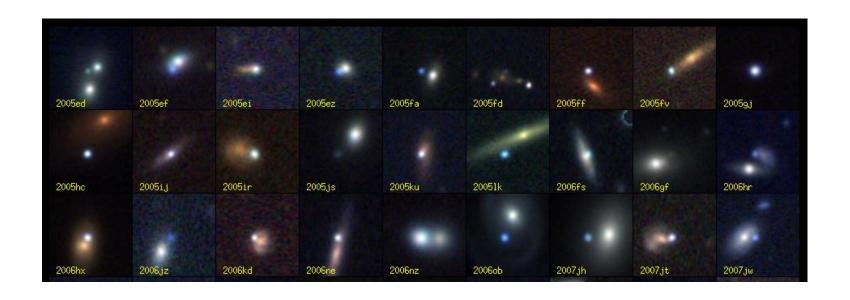


## Why is it hard?

- Clustering in two dimensions looks easy
- Clustering small amounts of data looks easy
- And in most cases, looks are not deceiving
- Many applications involve not 2, but 10 or 10,000 dimensions
- High-dimensional spaces look different: Almost all pairs of points are at about the same distance

## **Clustering Problem: Galaxies**

- A catalog of 2 billion "sky objects" represents objects by their radiation in 7 dimensions (frequency bands)
- Problem: Cluster into similar objects, e.g., galaxies, nearby stars, quasars, moons, belts, clouds, etc.
- Sloan Digital Sky Survey



## **Clustering Problem: Music CDs**

- Intuitively: Music divides into categories, and customers prefer a few categories
  - But what are categories really?
- Represent a CD by a set of customers who bought it:

 Similar CDs have similar sets of customers, and viceversa

## **Clustering Problem: Music CDs**

#### **Space of all CDs:**

- Think of a space with one dim. for each customer
  - Values in a dimension may be 0 or 1 only
  - A CD is a point in this space  $(x_1, x_2,..., x_k)$ , where  $x_i = 1$  iff the i th customer bought the CD
- For Amazon, the dimension is tens of millions
- Task: Find clusters of similar CDs

## Clustering Problem: Documents

### **Finding topics:**

- Represent a document by a vector  $(x_1, x_2,..., x_k)$ , where  $x_i = 1$  iff the  $i^{th}$  word (in some order) appears in the document
  - It actually doesn't matter if *k* is infinite; i.e., we don't limit the set of words
- Documents with similar sets of words may be about the same topic

## Clustering Problem: Topics (Dual)

#### **Finding topics:**

- Represent a document by a vector  $(x_1, x_2,..., x_k)$ , where  $x_i = 1$  iff the  $i^{th}$  word (in some order) appears in the document
  - It actually doesn't matter if *k* is infinite; i.e., we don't limit the set of words
- Topic is a set of similar documents

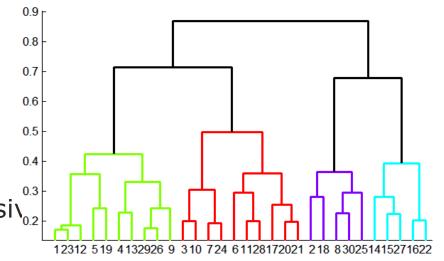
## Cosine, Jaccard, and Euclidean

- As with CDs we have a choice when we think of documents as sets of words:
  - Sets as vectors: Measure similarity by the cosine distance
  - Sets as sets: Measure similarity by the Jaccard distance
  - Sets as points: Measure similarity by Euclidean distance

## Overview: Methods of Clustering

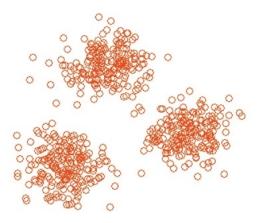
#### • Hierarchical:

- **Agglomerative** (bottom up):
  - ◆Initially, each point is a cluster
  - ◆ Repeatedly combine the two "nearest" clusters into one
- **Divisive** (top down):
  - ◆Start with one cluster and recursiv<sub>0.2</sub>



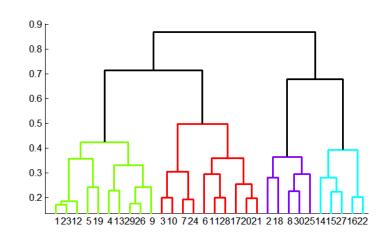
## Point assignment:

- Maintain a set of clusters
- Points belong to "nearest" cluster

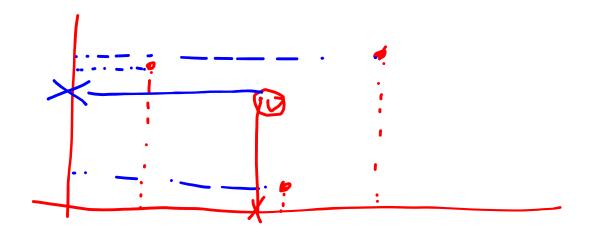


## **Hierarchical Clustering**

Key operation:
 Repeatedly combine
 two nearest clusters



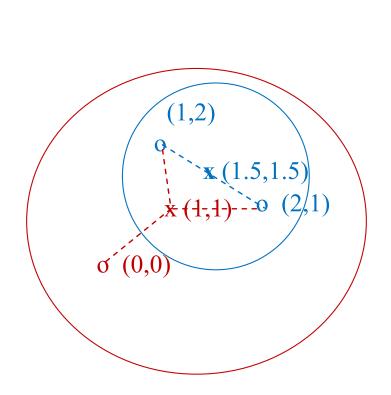
- Three important questions:
  - 1) How do you represent a cluster of more than one point?
  - 2) How do you determine the "nearness" of clusters?
  - 3) When to stop combining clusters?

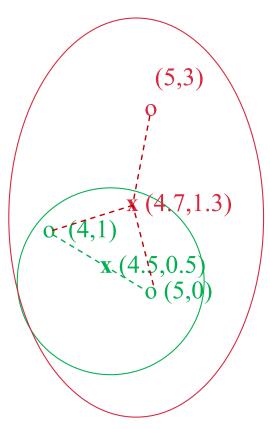


## **Hierarchical Clustering**

- Key operation: Repeatedly combine two nearest clusters
- (1) How to represent a cluster of many points?
  - **Key problem:** As you merge clusters, how do you represent the "location" of each cluster, to tell which pair of clusters is closest?
- Euclidean case: each cluster has a centroid = average of its (data)points
- (2) How to determine "nearness" of clusters?
  - Measure cluster distances by distances of centroids

## **Example: Hierarchical clustering**

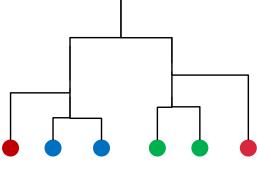




#### Data:

o ... data point

x ... centroid



**Dendrogram** 

#### And in the Non-Euclidean Case?

#### What about the Non-Euclidean case?

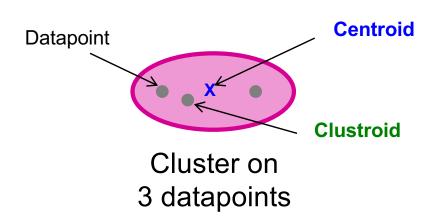
- The only "locations" we can talk about are the points themselves
  - i.e., there is no "average" of two points

#### Approach 1:

- (1) How to represent a cluster of many points? clustroid = (data)point "closest" to other points
- (2) How do you determine the "nearness" of clusters? Treat clustroid as if it were centroid, when computing inter-cluster distances

### "Closest" Point?

- (1) How to represent a cluster of many points?
   clustroid = point "closest" to other points
- Possible meanings of "closest":
  - Smallest maximum distance to other points
  - Smallest average distance to other points
  - Smallest sum of squares of distances to other points
    - ◆For distance metric **d** clustroid **c** of cluster **C** is:



$$\min_{c} \sum_{x \in C} d(x, c)^2$$

**Centroid** is the avg. of all (data)points in the cluster. This means centroid is an "artificial" point.

Clustroid is an existing (data)point that is "closest" to all other points in the cluster.

## **Defining "Nearness" of Clusters**

- (2) How do you determine the "nearness" of clusters?
  - Approach 2: Intercluster distance = minimum of the distances between any two points, one from each cluster
  - Approach 3: Pick a notion of "cohesion" of clusters, e.g., maximum distance from the clustroid
    - ◆Merge clusters whose *union* is most cohesive

## Cohesion

- Approach 3.1: Use the diameter of the merged cluster = maximum distance between points in the cluster
- Approach 3.2: Use the average distance between points in the cluster
- Approach 3.3: Use a density-based approach
  - Take the diameter or avg. distance, e.g., and divide by the number of points in the cluster

#### **Termination Conditions**

- When do you stop combining clusters?
- Approach 1: Pick k up front and stop when you have k
  - makes sense if we know that the data falls into k categories
- Approach 2: Stop when the next merge will create a cluster with low cohersion
  - i.e. bad cluster

## Implementation

- Naïve implementation of hierarchical clustering:
  - At each step, compute pairwise distances between all pairs of clusters, then merge
  - $\bigcirc (N^3)$
- Careful implementation using priority queue can reduce time to O(N<sup>2</sup> log N)
  - Still too expensive for really big datasets that do not fit in memory

# k-means clustering

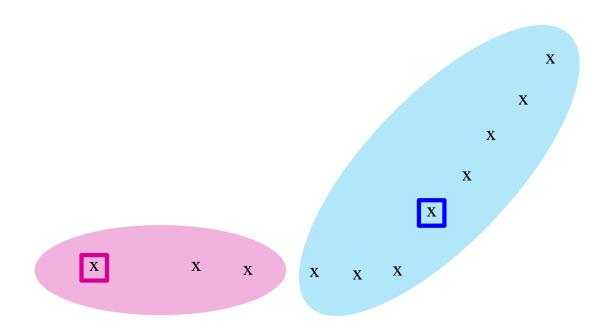
## *k*–means Algorithm(s)

- Assumes Euclidean space/distance
- Start by picking k, the number of clusters
- Initialize clusters by picking one point per cluster
  - **Example:** Pick one point at random, then **k-1** other points, each as far away as possible from the previous points

## **Populating Clusters**

- 1) For each point, place it in the cluster whose current centroid it is nearest
- 2) After all points are assigned, update the locations of centroids of the k clusters
- 3) Reassign all points to their closest centroid
  - Sometimes moves points between clusters
- Repeat 2 and 3 until convergence
  - Convergence: Points don't move between clusters and centroids stabilize

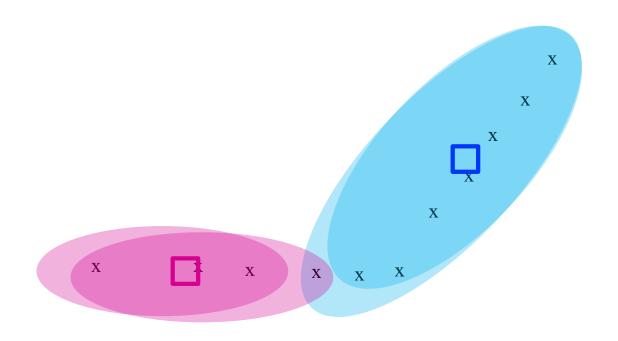
# **Example: Assigning Clusters**



x ... data point ... centroid

**Clusters after round 1** 

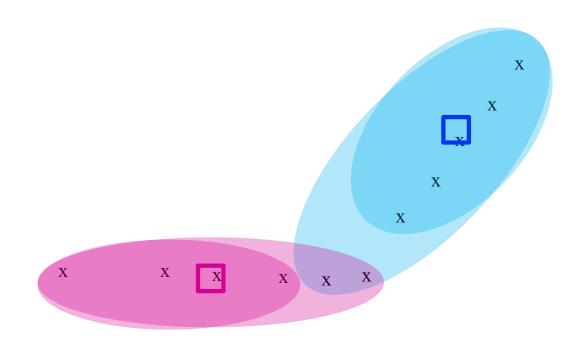
# **Example: Assigning Clusters**



x ... data point ... centroid

**Clusters after round 2** 

## **Example: Assigning Clusters**

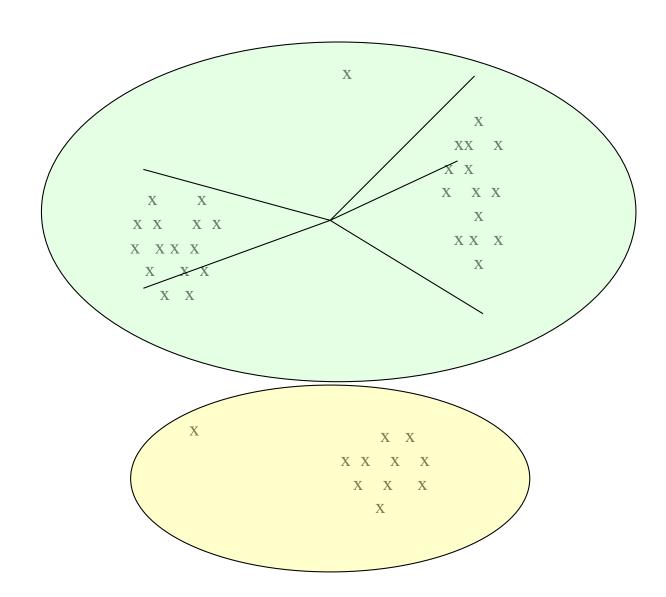


x ... data point ... centroid

Clusters at the end

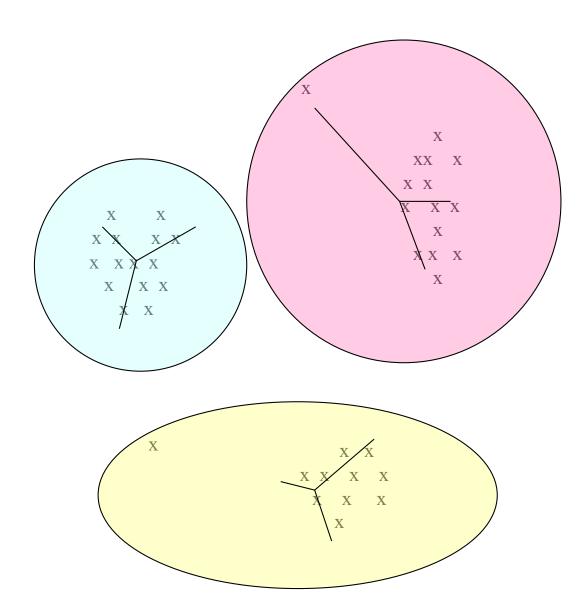
## Example: Picking k

Too few; many long distances to centroid.




# Example: Picking k

Just right; distances rather short.



## Example: Picking k


Too many; little improvement in average distance.



## Getting the k right

#### How to select *k*?

- Try different k, looking at the change in the average distance to centroid as k increases
- Average falls rapidly until right k, then changes little



## Picking the initial k points

#### Approach 1: Sampling

- Cluster a sample of the data using hierarchical clustering, to obtain k clusters
- Pick a point from each cluster (e.g. point closest to centroid)
- Sample fits in main memory

#### Approach 2: Pick "dispersed" set of points

- Pick first point at random
- Pick the next point to be the one whose minimum distance from the selected points is as large as possible
- Repeat until we have k points

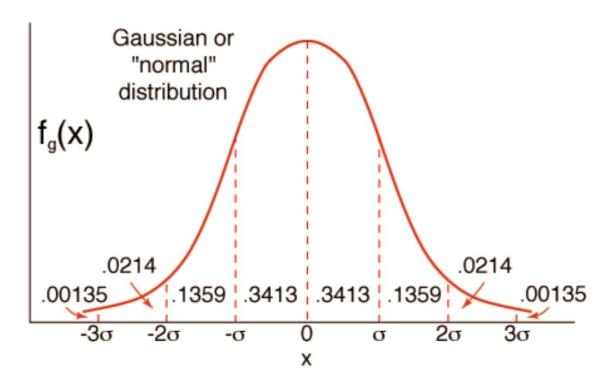
# The BFR Algorithm

Extension of k-means to larger data

# **BFR Algorithm**

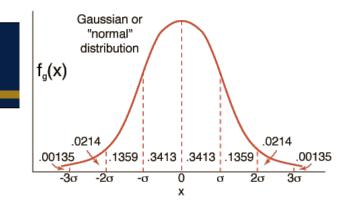
- BFR [Bradley-Fayyad-Reina] is a variant of kmeans for very large (disk-resident) data sets
- Assumes each cluster is normally distributed around a centroid in Euclidean space

# **Normal Distribution**

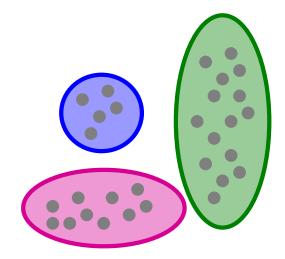


- Can quantify the likelihood of finding a point in the cluster at a given distance from the centroid along each dimension
- Standard deviations in different dimenions may vary

## **BFR Algorithm**



- BFR [Bradley-Fayyad-Reina] is a variant of k-means designed to handle very large (disk-resident) data sets
- Assumes that clusters are normally distributed around a centroid in a Euclidean space
  - Standard deviations in different dimensions may vary
    - ◆ Clusters are axis-aligned ellipses
- Efficient way to summarize clusters
   (memory required: O(clusters) and not O(data))



### **BFR Algorithm**

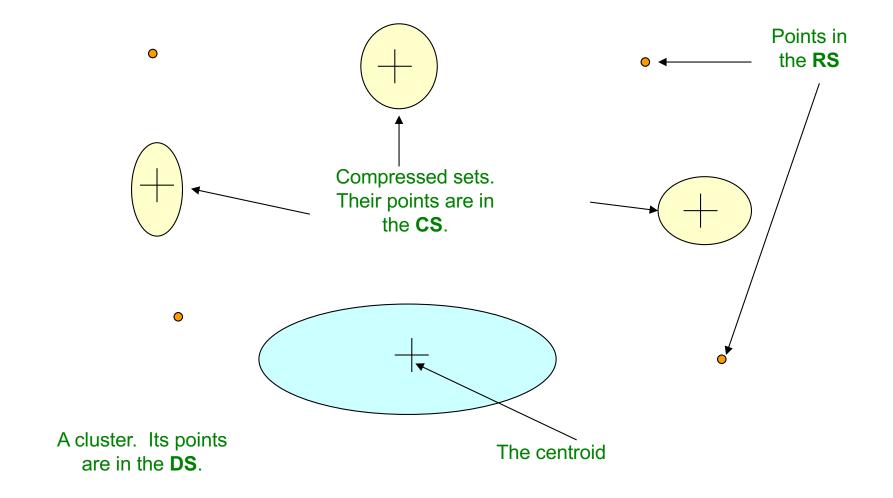
- Points are read from disk one main-memory-full at a time
- Most points from previous memory loads are summarized by simple statistics
- To begin, from the initial load we select the initial k
  centroids by some sensible approach:
  - Take **k** random points
  - Take a small random sample and cluster optimally
  - Take a sample; pick a random point, and then k-1 more points, each as far from the previously selected points as possible

#### Three Classes of Points

#### 3 sets of points which we keep track of:

- Discard set (DS):
  - Points close enough to a centroid to be summarized
- Compression set (CS):
  - Groups of points that are close together but not close to any existing centroid
  - These points are summarized, but not assigned to a cluster
- Retained set (RS):
  - Isolated points waiting to be assigned to a compression set

#### **BFR: "Galaxies" Picture**



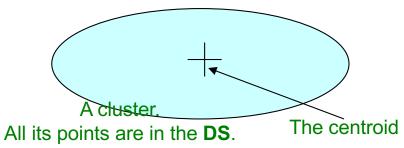
Discard set (DS): Close enough to a centroid to be summarized Compression set (CS): Summarized, but not assigned to a cluster Retained set (RS): Isolated points

43

## Summarizing Sets of Points

#### For each cluster, the discard set (DS) is <u>summarized</u> by:

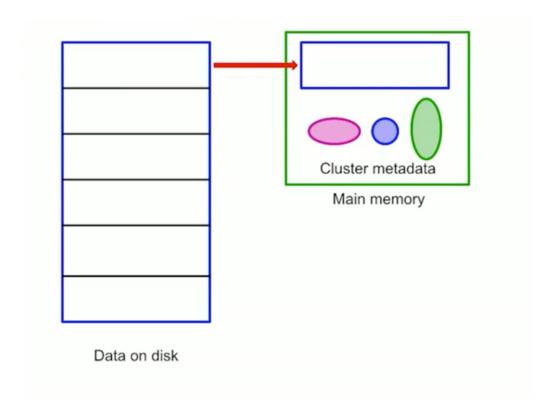
- The number of points, N
- The vector *SUM*, whose *I*<sup>th</sup> component is the sum of the coordinates of the points in the *I*<sup>th</sup> dimension
- The vector **SUMSQ**:  $t^h$  component = sum of squares of coordinates in  $t^h$  dimension



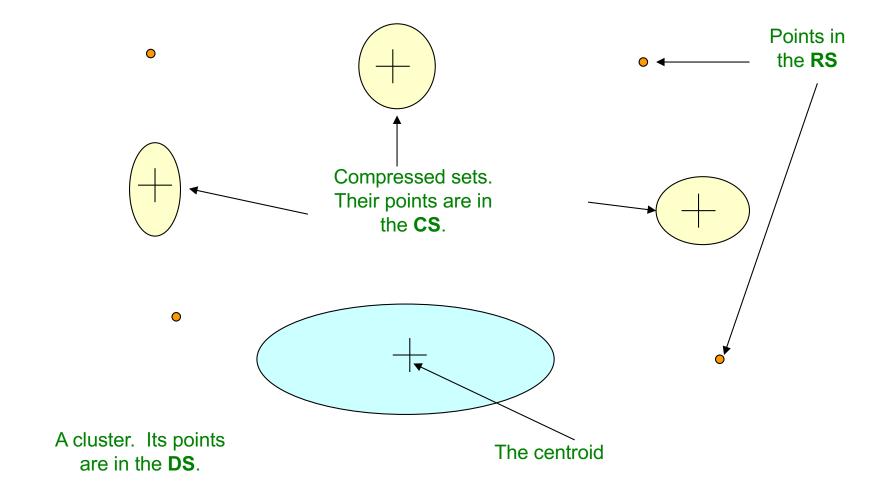
## **Summarizing Sets of Points**

- 2d + 1 values represent any size cluster
  - d = number of dimensions
- Average in each dimension (the centroid)
   can be calculated as SUM; / N
  - SUM<sub>i</sub> = i<sup>th</sup> component of SUM
- Variance of a cluster's discard set in dimension i is: (SUMSQ<sub>i</sub> / N) – (SUM<sub>i</sub> / N)<sup>2</sup>
  - And standard deviation is the square root of that
- Next step: Actual clustering

# **BFR Overview**



#### **BFR: "Galaxies" Picture**



Discard set (DS): Close enough to a centroid to be summarized Compression set (CS): Summarized, but not assigned to a cluster Retained set (RS): Isolated points

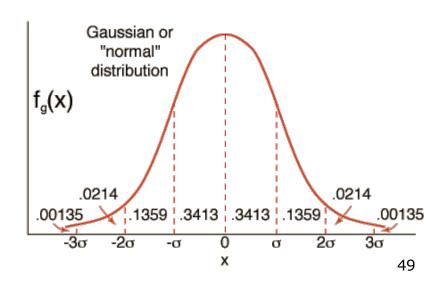
47

#### A Few Details...

- Q1) How do we decide if a point is "close enough" to a cluster that we will add the point to that cluster?
- Q2) How do we decide whether two compressed sets (CS) deserve to be combined into one?
  - Compute the variance of the combined subcluster

### **How Close is Close Enough?**

- Q1) We need a way to decide whether to put a new point into a cluster (and discard)
- BFR suggests two ways:
  - The Mahalanobis distance is less than a threshold
  - High likelihood of the point belonging to currently nearest centroid



#### **Mahalanobis Distance**

#### Normalized Euclidean distance from centroid

- For point  $(x_1, ..., x_d)$  and centroid  $(c_1, ..., c_d)$ 
  - 1. Normalize in each dimension:  $y_i = (x_i c_i) / \sigma_i$
  - 2. Take sum of the squares of the  $y_i$
  - 3. Take the square root

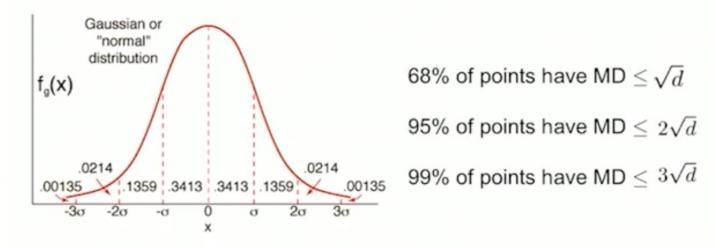
$$d(x,c) = \sqrt{\sum_{i=1}^{d} \left(\frac{x_i - c_i}{\sigma_i}\right)^2}$$

measures the # of standard deviations from centroid across a dimension

 $\sigma_i$  ... standard deviation of points in the cluster in the  $i^{th}$  dimension

## **Mahalanobis Acceptance Criterion**

- Suppose point P is one standard deviation away from centroid in each dimension
  - Each  $y_i = 1$  and so the MD of P is  $\sqrt{d}$



Accept point P into cluster C if its MD from cluster centroid is less than a threshold e.g.,  $3\sqrt{d}$ 

#### **Should 2 Compressed Set be Combined?**

#### Q2) Should 2 CS subclusters be combined?

- Compute the variance of the combined subcluster
  - N, SUM, and SUMSQ allow us to make that calculation quickly
- Combine if the combined variance is below some threshold





**Variance** = Square of Standard Deviation

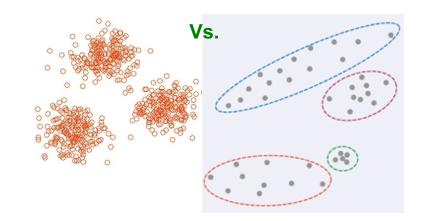
# The CURE Algorithm

# Extension of *k*-means to clusters of arbitrary shapes

## The CURE Algorithm

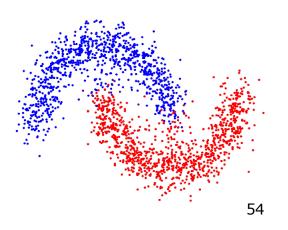
#### Problem with BFR/k-means:

- Assumes clusters are normally distributed in each dimension
- And axes are fixed ellipses at an angle are *not OK*

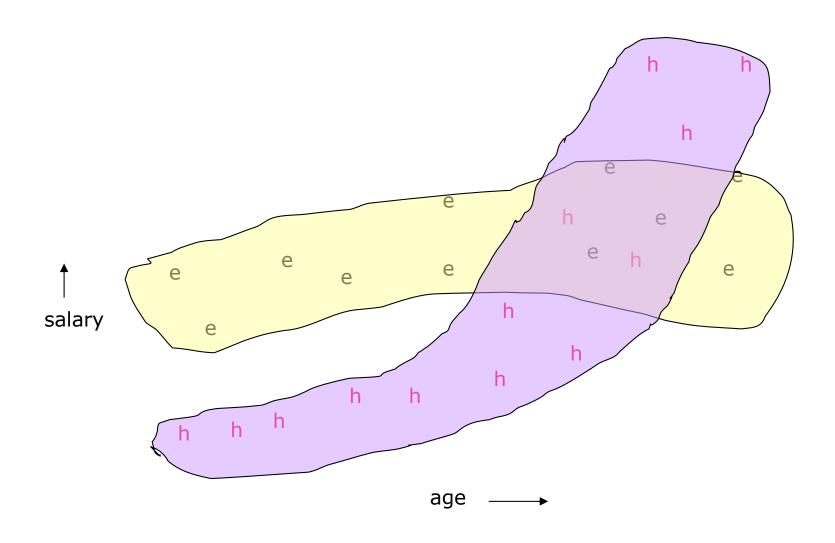


#### CURE (Clustering Using REpresentatives):

- Assumes a Euclidean distance
- Allows clusters to assume any shape
- Uses a collection of representative points to represent clusters



# **Example: Professor's Salaries**

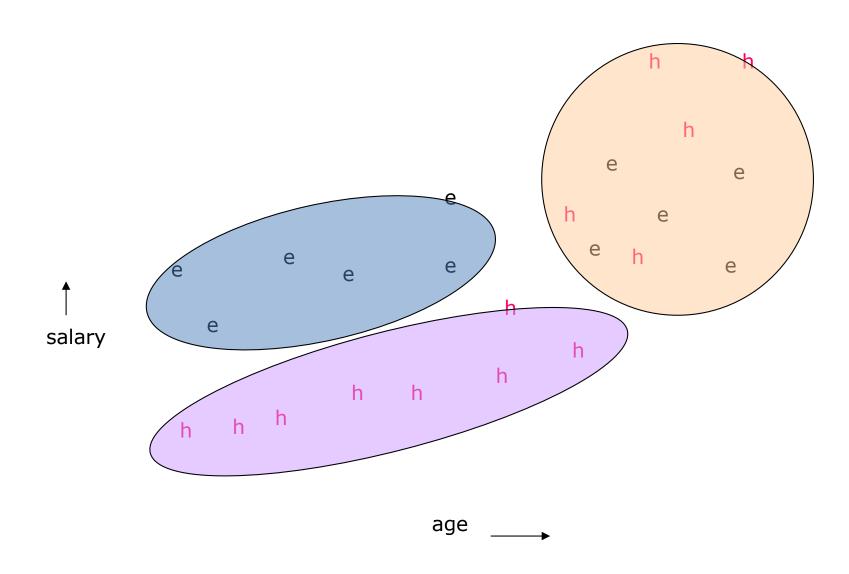


## **Starting CURE**

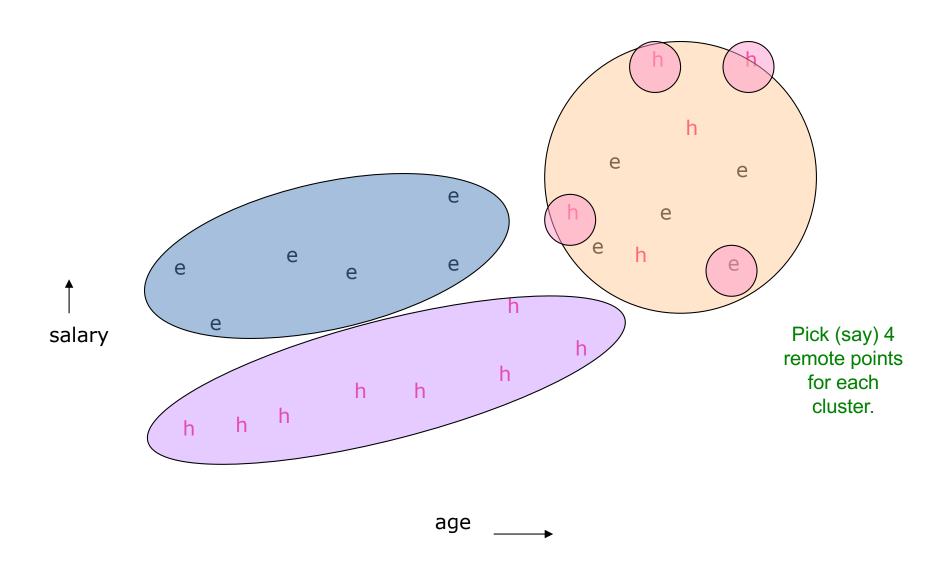
#### 2 Pass algorithm. Pass 1:

- 0) Pick a random sample of points that fit in main memory
- 1) Initial clusters:
  - Cluster these points hierarchically group nearest points/clusters
- 2) Pick representative points:
  - For each cluster, pick a sample of points, as dispersed as possible
  - From the sample, pick representatives by moving them (say) 20% toward the centroid of the cluster

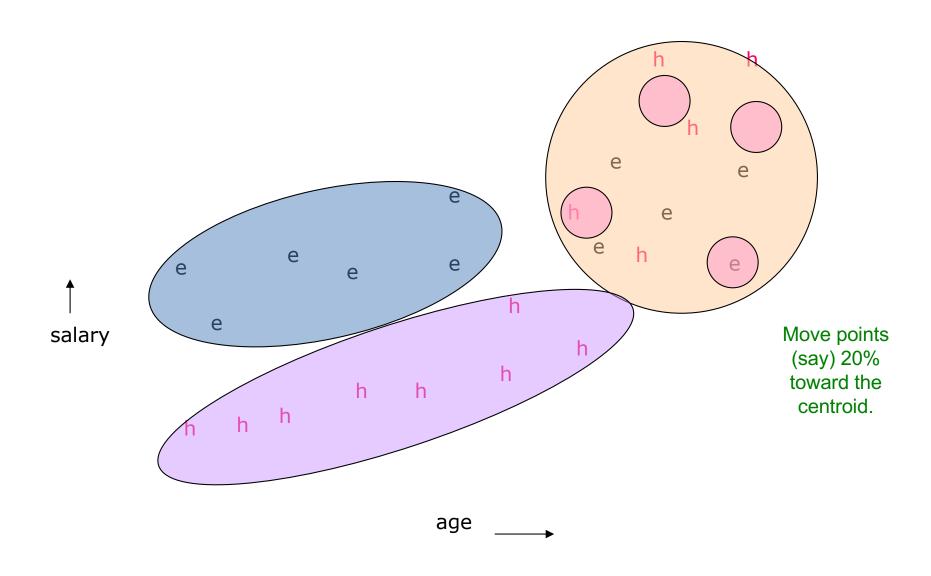
# **Example: Initial Clusters**



# **Example: Pick Dispersed Points**



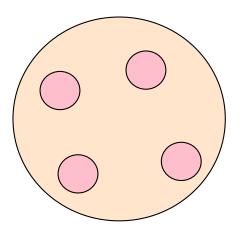
# **Example: Pick Dispersed Points**



## **Finishing CURE**

#### **Pass 2:**

- Now, rescan the whole dataset and visit each point p in the data set
- Place it in the "closest cluster"
  - Normal definition of "closest": Find the closest representative to **p** and assign it to representative's cluster



p

## Summary

- Clustering: Given a set of points, with a notion of distance between points, group the points into some number of clusters
- Algorithms:
  - Agglomerative hierarchical clustering:
    - ◆Centroid and clustroid
  - *k*-means:
    - lacktriangle Initialization, picking k
  - BFR
  - CURE



#### **DBSCAN**

<u>Density-based Clustering</u> locates regions of high density that are separated from one another by regions of low density.

- Density = number of points within a specified radius (Eps)
- DBSCAN is a density-based algorithm.
  - A point is a core point if it has more than a specified number of points (MinPts) within Eps
    - ◆These are points that are at the interior of a cluster
  - A border point has fewer than MinPts within Eps, but is in the neighborhood of a core point

#### **DBSCAN**

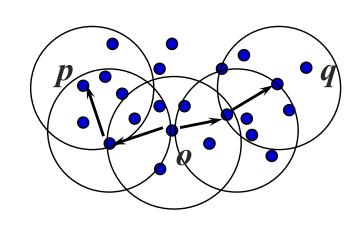
- A noise point is any point that is not a core point or a border point.
- Any two core points are close enough— within a distance Eps of one another – are put in the same cluster
- Any border point that is close enough to a core point is put in the same cluster as the core point
- Noise points are discarded

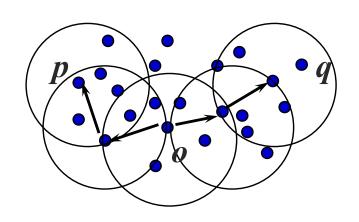
### **DBSCAN: The Algorithm**

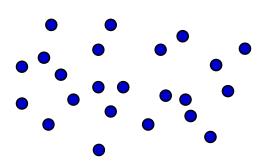
- select a point p
- Retrieve all points density-reachable from p wrt  $\epsilon$  and MinPts.
- If p is a core point, a cluster is formed.
- If **p** is a border point, no points are density-reachable from **p** and DBSCAN visits the next point of the database.
- Continue the process until all of the points have been processed.

Result is independent of the order of processing the points

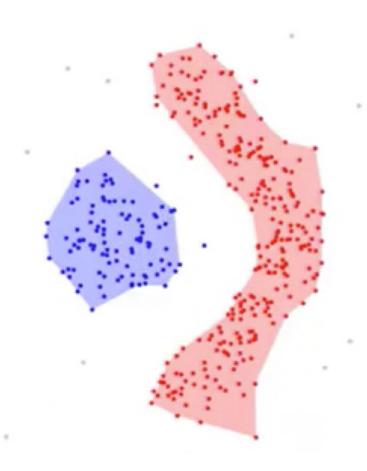
# An example







# **DB Scan Clusters**



## **DBSCAN** Disadvantage: Irregular Density

